

TRACKD

Product Requirements Document

Scraper

Job Tracker

Analytics

AI Insights

Self-Growth

PREPARED BY
PRODUCT
VERSION
DATE
CLASSIFICATION

VybzTech Inc.
TRACKD — AI Job Application Tracker
1.0 | Draft for Stakeholder Review
July 2026
Confidential

Table of Contents

- 1** Overview — What TRACKD Is and Why It Exists
- 2** The Problem We Are Solving
- 3** Scope — What We Are Building, and What We Are Not
- 4** The Scraper — Two Ways In
 - 4.1 Method One: Smart Paste
 - 4.2 Method Two: URL Ingestion
 - 4.3 The Ingestion Side Panel
- 5** Job Tracking — The Dashboard
 - 5.1 Kanban View
 - 5.2 Calendar View
 - 5.3 Table View
 - 5.4 Status Lifecycle
- 6** Analytics
 - 6.1 KPI Cards
 - 6.2 Chart Suite
- 7** AI Insights
 - 7.1 Match Score
 - 7.2 ATS Keyword Analysis
 - 7.3 Resume Bullet Suggestions
- 8** Self-Improvement — The Empowerment Layer
 - 8.1 The Resume Canvas
 - 8.2 Cover Letter Generator
 - 8.3 Growth Loop
- 9** User Stories & Acceptance Criteria
- 10** Non-Functional Requirements
- 11** Risks & Mitigations
- 12** Phased Roadmap

1. Overview — What TRACKD Is and Why It Exists

Job hunting at scale is a management problem disguised as a motivation problem. Most people think they are struggling because they are not applying to enough roles or not writing compelling enough cover letters. The real issue is usually operational: their process is a mess, their data lives in five different places, and they have no honest view of what is working and what is not.

TRACKD is the command centre that fixes the operational problem. It captures every job opportunity from a URL or a paste, structures that raw text into clean data, and puts it into a live pipeline you can see, sort, and act on. An AI layer sits on top: it scores your resume against each role, flags what is missing, rewrites your bullet points for the specific job, and drafts a cover letter in under thirty seconds. The analytics page tells you how your search is actually going — not how you feel it is going.

There are five things TRACKD does. The scraper gets job data in with zero manual typing. The dashboard gives you a pipeline you can see and manage. The analytics page surfaces the numbers. The AI insight engine tells you where you are strong and where you are exposed for each specific role. And the self-improvement layer gives you the tools to fix what the AI finds — editable resume canvas, one-click cover letters, tailored to the job.

This document is the single authoritative reference for what TRACKD must do in its first production release. It covers the five pillars above in sequence: what the feature is, why it works the way it does, and the exact behaviour the product must deliver.

2. The Problem We Are Solving

It is worth being honest about what is actually broken, because every design decision in this document traces back to it.

The job hunt is high-volume and manual

A competitive job seeker applies to between 30 and 150 roles over a typical search cycle. Each application requires copying the same personal information, tailoring the same resume, and writing a cover letter that is structurally identical to the previous twenty. The sheer repetition is not just exhausting — it produces worse output. After the fifteenth near-identical cover letter, quality declines. After the thirtieth application, people stop tracking at all and start applying blindly.

The data is nowhere and everywhere

The average job seeker manages their pipeline across a personal spreadsheet they stopped updating three weeks ago, a folder of screenshots, a collection of open browser tabs, and a notebook. There is no single place to see all applications, their statuses, and what needs to happen next. Opportunities fall through because there was no system to catch them.

There is no performance data

People finish a job search without understanding why they got interviews for some roles and silence for others. They do not know their response rate, their average time to first response, which platforms gave them the best return, or whether their resume is passing ATS screening. Without data, the only feedback loop is intuition — and intuition says 'try harder', which is not useful advice.

Application materials are generic

The cover letter problem is not that people cannot write. It is that writing a genuinely tailored letter for every role is time-prohibitive. The result is a generic letter that reads like a generic letter, which performs worse in both human and automated screening.

The design principle that follows from all of this:

Every TRACKD feature must either reduce friction on the input side (getting data in), improve visibility on the management side (understanding the pipeline), or improve quality on the output side (making each application stronger). Features that do not fit one of these three categories do not belong in the first release.

3. Scope — What We Are Building, and What We Are Not

Scope discipline matters at this stage. The temptation is always to add features that sound useful but push the first launch back by months. For version one, here is the line.

Feature	Phase
Scraper — two-method ingestion (paste + URL)	Phase 1
Dashboard — Kanban, Calendar, and Table views	Phase 1
Status tracking with full history log	Phase 1
Analytics page with five core charts and KPI cards	Phase 1
AI match score, missing skills, ATS keyword risk	Phase 1
Resume canvas editor with AI suggestions	Phase 1
Cover letter generator (three tone options)	Phase 1
User profile and onboarding (resume extraction)	Phase 1
Free and Pro tier access gates	Phase 1
Browser extension — one-click save from any job board	Phase 2
Email webhook — auto-status from recruiter emails	Phase 2
Interview prep — AI practice questions per role	Phase 3
Predictive analytics — response probability models	Phase 3

Out of scope for Phase 1

- Native mobile app (iOS or Android)
- Team or recruiter-facing multi-seat workspace
- Direct job application submission on behalf of the user
- Third-party calendar sync (Google Calendar, Outlook)
- Automatic job board scraping without user-provided input

4. The Scraper — Two Ways In

The single biggest source of drop-off in any tracking tool is the moment the user has to type. The moment someone has to manually enter a company name, a job title, and a salary range from a job posting they already have open in another tab, a portion of them will decide it is not worth it. TRACKD removes that moment entirely.

There are two ingestion methods in Phase 1. Both end at the same place: a structured job data object in Firestore that populates the dashboard. The difference is the starting point.

4.1 Method One — Smart Paste

The user copies the job description from wherever they found it — LinkedIn, Indeed, a company careers page, a PDF, a recruiter's email — and pastes it into the Smart Input Area on the Ingestion page. They click Parse. That is it.

Behind the scenes, the Go backend receives the raw text, strips any residual HTML or formatting noise, and submits it to Gemini with a strict JSON schema that defines every field the system needs: company name, role title, department, seniority level, work mode, minimum and maximum compensation, currency, equity, benefits, technical stack, ATS keywords, and application deadline. Gemini returns the structured JSON. The backend validates it. The side panel opens.

Smart Paste is the primary method for Phase 1 because it works everywhere a user can copy text. It does not depend on any job board's DOM structure staying the same. It does not require a browser extension to be installed. It works from a recruiter email, a PDF, a forwarded message, or a screenshot-to-text paste.

4.2 Method Two — URL Ingestion

The user pastes a URL from a public job listing. The Go backend fetches the page content, uses goquery to strip navigation, headers, footers, and irrelevant elements, and extracts the clean body text. From that point the flow is identical to Smart Paste: Gemini receives the cleaned text, returns structured JSON, and the side panel opens.

URL ingestion is a convenience method for public listings that the user has not already copied. It fails gracefully: if the backend encounters a login wall, a redirect, or a 403 response, it returns a clear error and prompts the user to switch to Smart Paste. The system never pretends a parse succeeded when it did not.

Why not auto-scrape job boards directly?

Scraping job boards at scale requires maintaining scrapers for dozens of platforms that change their HTML structure without notice. It also crosses into legally grey territory with some platforms' terms of service. User-initiated URL ingestion — where the user deliberately provides a link they want to track — is a fundamentally different act, both technically and legally. We do it right in Phase 1, and we build the browser extension in Phase 2 for users who want one-click saving.

4.3 The Ingestion Side Panel

Every parse, regardless of method, produces a side panel that slides in from the right before anything is written to the database. This is a deliberate design choice. The AI will occasionally get something wrong — a salary range it could not find, a work mode it misread, a job title with an unusual format. The user must be able to see and correct the parsed output before it becomes permanent.

Field Group	Fields Extracted
Role	Job title, department, seniority level, work mode (remote/hybrid/onsite)
Company	Company name, industry, approximate size
Compensation	Salary minimum, salary maximum, currency, pay type, equity, key benefits
Timeline	Application deadline (if stated)
Tech Requirements	Technology stack, required tools, ATS-relevant keywords

Each field carries a confidence indicator: green for high confidence, amber for medium, and red for low or absent. Low-confidence fields are visually prominent so the user knows exactly where to look. The Commit button is disabled until at minimum the company name and role title are populated — either by the AI or by the user. Nothing goes to Firestore that is not at least minimally complete.

5. Job Tracking — The Dashboard

Once a job is committed, it lands in the dashboard. The dashboard is the central workspace. It is where the user manages their pipeline day to day: updating statuses, reviewing what is active, setting priorities, and opening jobs to work on their applications. It offers three views of the same underlying data. The user switches between them freely and the data is always in sync.

5.1 Kanban View

The default view. Five columns represent the five possible states of any application. Cards move between columns by drag-and-drop. Every move is a status update in Firestore. Every status update is logged in the job's history. The user gets a real-time visual representation of exactly where every opportunity stands.

Column Label	Status Value	Meaning
Saved	Bookmarked	The role has been captured but the user has not applied yet
Applied	Applied	Application has been submitted
Interviewing	Interviewing	At least one interview has been confirmed or completed
Offer Received	Offer	A formal offer has been made
Rejected / Closed	Rejected	Application has been declined or the role has been filled

Each card shows the company name (with auto-fetched logo where available), the job title, the date it was logged, the source, and — for Pro users — the AI match score as a small circular badge. Right-clicking or opening the card's context menu gives access to View Pro Page, Archive, and Delete.

5.2 Calendar View

Jobs appear on their applied date. A second marker appears on any application deadline date that was captured during ingestion. The user can switch between month, week, and day views. Clicking any event opens a quick-view modal with core job details and a link to the Pro page. The calendar view solves the deadline blindspot — the user can see at a glance what is coming up and plan accordingly.

5.3 Table View

A dense, sortable, filterable grid for users who prefer working with data in rows. Every column is sortable. Status and source are multi-select filterable. The status dropdown is editable inline — the

user does not need to open the Pro page to update where an application stands. Row checkboxes enable bulk archive or bulk delete. Clicking a row navigates to the Pro page.

Column	Behaviour
Company	Sortable
Role Title	Sortable
Status	Sortable + Multi-select filter
Match Score	Sortable — Pro users only
Applied Date	Sortable + Date range filter
Deadline	Sortable + Date range filter
Source	Multi-select filter
Actions	View Pro Page, Archive, Delete

5.4 Status Lifecycle

Status transitions follow a defined matrix. Not every transition is permitted — you cannot move directly from Saved to Offer, for example. The allowed paths reflect how real hiring processes work. Every transition is recorded in the statusHistory log with a timestamp, and this history is what powers the time-in-stage analytics.

From Status	Allowed Next Statuses
Saved (bookmarked)	Applied, Rejected
Applied	Interviewing, Rejected
Interviewing	Offer Received, Rejected
Offer Received	Terminal — no further transition
Rejected	Terminal — no further transition

6. Analytics

The analytics page exists for one reason: to give users honest data about their own search. Not encouragement, not gamification — data. How many applications have gone out. What percentage are getting responses. How long it is taking to hear back. Which platforms are producing the best return. Where applications are getting stuck in the funnel.

Most people finishing a job search could not answer any of those questions with numbers. TRACKD makes every one of them answerable, in real time, with no manual calculation required.

6.1 KPI Cards

Four headline numbers sit at the top of the analytics page. They update within five seconds of any status change across the dashboard.

Metric	Calculation	Why It Matters
Total Applications	Count of all non-archived tracked jobs	Gives a sense of volume and effort invested
Response Rate	$(\text{Interviewing} + \text{Offer}) / \text{Applied} \times 100\%$	Honest measure of how well the pipeline is converting
Offer Rate	$\text{Offer} / \text{Applied} \times 100\%$	The number that matters most at the end
Avg Time to Response	Mean days from Applied to first Interviewing status entry	Identifies whether the delay is in applications or in interviews

6.2 Chart Suite

Five charts sit below the KPI cards. A date range filter at the top — Last 30 Days, Last 90 Days, All Time — applies to all charts simultaneously.

Chart	Type	What It Shows
Applications Over Time	Area line chart	Weekly or monthly count of applications submitted — shows effort and pacing
Pipeline Breakdown	Donut chart	Current distribution of all applications across the five status values
Applications by Source	Horizontal bar chart	Count per source (LinkedIn, Indeed, referral, extension, paste) — shows where to spend time
Top Roles Applied To	Ranked list	Frequency of each job title in tracked applications

Chart	Type	What It Shows
Average Time in Stage	Stat cards (one per status)	Mean days spent in each status — surfaces where the process is slow

Note on data quality:

Analytics are only as good as the status update discipline of the user. The onboarding flow explains this. The dashboard empty-column prompts reinforce it. If the user stops updating statuses, the analytics become stale and misleading — and we would rather show an honest warning than a false picture.

7. AI Insights

The AI insight layer is what separates TRACKD from a well-designed spreadsheet. It takes the user's profile — their resume, their skills, their career trajectory — and cross-references it against each specific job they have tracked. The output is not a generic score. It is a per-job diagnosis: here is where you are strong, here is where you are exposed, and here are the exact words you are missing that this company's ATS is likely looking for.

The AI insight panel lives on the Pro page, accessible from any job card or table row. Free-tier users see a preview of the panel with an upgrade prompt. Pro users see the full output.

7.1 Match Score

A 0–100 percentage shown as a circular gauge at the top of the AI insights panel. It represents how well the user's current profile aligns with the stated requirements of the specific job description. The colour of the gauge changes with the score: green at 70 and above, amber between 40 and 69, red below 40.

The score is calculated by Gemini comparing the user's extracted skills, experience titles, and seniority level against the job's required and preferred qualifications. It is a directional signal, not a guarantee. We are honest with the user about this: the score is a tool, not a verdict.

7.2 ATS Keyword Analysis

Applicant Tracking Systems scan resumes for specific terms before a human ever reads them. A resume that is a perfect fit for a role but uses different vocabulary will fail at this stage. The ATS keyword analysis component shows:

- A risk level badge — Low, Medium, or High — based on how many required keywords are absent from the user's resume
- An expandable list of the specific keywords that are flagged as missing or underrepresented
- Where each keyword appeared in the job description (required versus preferred section)

This output is directly actionable: the user opens the resume editor, accepts the AI suggestion that adds the flagged keyword in context, and the risk level updates. The loop between insight and action is closed within the same page.

7.3 Resume Bullet Suggestions

The most opinionated part of the AI layer. Gemini reviews the user's existing resume experience bullets and rewrites them for the specific job: stronger action verbs where the current ones are weak, quantified impact where the user has left it vague, and terminology that matches the job description without misrepresenting the user's actual experience.

Each suggestion appears as a chip above the relevant bullet in the content editor. The user can accept it — replacing the original with the suggestion — or dismiss it. Every accepted suggestion is reflected immediately on the live resume canvas to the left. Nothing is automatic. The user is always in control of what goes into their resume.

AI suggestion quality depends on the completeness of the user's profile. A profile with a fully uploaded resume, confirmed extracted experience, and set career preferences will produce significantly more relevant suggestions than a sparse profile. The onboarding flow, the profile completeness score, and the nudge prompts all exist to push users toward a complete profile for this reason.

8. Self-Improvement — The Empowerment Layer

The AI insights panel tells you what is wrong. The self-improvement layer gives you the tools to fix it, without leaving TRACKD, without opening a separate document editor, and without copying and pasting between five windows.

This is the part of the product that converts an insight into an outcome. A match score of 52% is only useful if the user can do something about it in the next five minutes. The resume canvas and the cover letter generator make that possible.

8.1 The Resume Canvas

A split-pane layout on the Pro page. The left pane renders a print-ready resume — the user's profile, formatted and live. The right pane is the content editor: tabbed across Summary, Experience, Skills, and Education. Every edit in the right pane updates the left pane in under 200 milliseconds. The user sees exactly what their resume will look like as they work on it.

AI suggestions appear inline in the content editor as highlighted chips. The user works through them — accepting the ones that fit, dismissing the ones that do not — and the resume on the left updates as they go. When they are done, they click Download. A server-side PDF is generated and downloaded to their device, named for the company and role: `FirstName_LastName_CompanyName_Resume.pdf`.

Why server-side PDF, not browser print-to-PDF?

Browser print-to-PDF produces inconsistent output depending on the browser, the OS, and the user's printer settings. A resume that looks correct in Chrome on macOS may have broken pagination in Firefox on Windows. Server-side rendering via a headless Chrome instance (chromedp) produces an identical PDF regardless of what the user is running. For a Pro feature that a user is submitting to a hiring manager, consistency is not optional.

8.2 Cover Letter Generator

The cover letter generator sits below the AI insights panel on the right pane of the Pro page. The user selects a tone, clicks Generate, and Gemini produces a full cover letter using the job description and the user's profile as context. The output appears in an editable text area.

Tone Option	Character
Professional	Formal, measured, focused on credentials and track record. Best for corporate and financial roles.

Tone Option	Character
Confident	Direct and assertive, focused on impact and initiative. Best for growth, product, and leadership roles.
Conversational	Warmer register, more personality, still professional. Best for startup, creative, and culture-forward roles.

The generated letter is editable. The user can change any line before downloading or copying it. A Regenerate option is available with a confirmation prompt if the user wants a fresh version. Download produces the same server-side PDF pipeline as the resume, formatted as a single-page letter.

8.3 The Growth Loop

The self-improvement layer closes a loop that no other tool in this category has built. The analytics page shows the user their response rate is low. The AI insights panel shows them that their ATS keyword risk is high across most of their recent applications. The resume canvas lets them fix it. They accept the AI suggestions, download the updated resume, apply again, and wait for the next data point.

Over time — across multiple applications, multiple iterations, multiple rounds of analytics — the user gets better at applying. Their match scores improve. Their response rate improves. The product has an active interest in their success, not just their continued engagement. That is what the empowerment layer is for.

Phase 3 extends this loop with interview preparation: AI-generated practice questions tailored to the specific company and role, based on the job description and what is known about the company's interview style. The foundation for this is the job data and profile data that already exists in the system by that point.

9. User Stories & Acceptance Criteria

ID	Feature	Story	Acceptance Criteria
US-01	Smart Paste	As a job seeker I want to paste a job description and have all fields auto-populated so I do not have to type anything	Given a valid job description paste, when I click Parse, the side panel opens within 8 seconds with all extractable fields populated
US-02	URL Parse	As a user I want to paste a job URL and have the system fetch and parse it	Given a valid public URL, when I click Parse, the backend fetches, cleans, and extracts the job data within 10 seconds
US-03	Field Edit	As a user I want to correct AI-extracted fields before committing	Given a populated side panel, every field is editable; changes persist to Firestore on commit
US-04	Kanban Drag	As a user I want to drag a job card to update its status	Given any job card, drag-drop to a new column updates the status in Firestore and appends to statusHistory within 2 seconds
US-05	Calendar	As a user I want to see my applied dates and deadlines on a calendar	Given jobs with dates, Calendar view shows each job on its applied date with a second marker on any deadline
US-06	Analytics	As a user I want charts showing my application activity	Given 5+ applications, Analytics shows all 5 charts and 4 KPI cards with correct values
US-07	Match Score	As a Pro user I want a match score for each job	Given a Pro page for any job, match score, missing skills, and ATS risk are displayed within 8 seconds of first load
US-08	Accept Suggestion	As a Pro user I want to accept an AI resume suggestion	Given a suggestion chip, clicking Accept updates the editor and canvas within 200ms
US-09	Resume Download	As a Pro user I want to download a tailored resume PDF	Given the Pro page, Download Resume produces a correctly formatted PDF within 5 seconds
US-10	Cover Letter	As a Pro user I want to generate a cover letter with my chosen tone	Given a tone selection, Generate produces a relevant letter in the output area within 8 seconds
US-11	Free Tier Gate	As a Free user accessing Pro features I want a clear upgrade path	Given a Free user on any Pro page, a ProGate modal is shown with feature list and upgrade CTA before any Pro content is visible

ID	Feature	Story	Acceptance Criteria
US-12	Resume Re-upload	As a user I want to update my profile by uploading a new resume	Given Profile page, uploading a new PDF triggers a new Gemini extraction and replaces the existing UserProfile on confirm

10. Non-Functional Requirements

Category	Requirement	Target
Performance	Ingestion API response (excl. Gemini)	< 3 seconds
Performance	Gemini extraction response time	< 8 seconds (P95)
Performance	Dashboard load (50 jobs)	< 3 seconds (P95)
Performance	Resume canvas edit-to-render	< 200ms
Performance	Analytics page load	< 3 seconds (P95)
Availability	Platform uptime	99.5% monthly SLA
Scalability	Concurrent users without degradation	1,000+ via Cloud Run auto-scaling
Security	Authentication	Firebase JWT validated on every API call
Security	Data encryption	Firestore at-rest encryption; HTTPS enforced
Security	OWASP Top 10	Validated before production launch
Accessibility	Standard	WCAG 2.1 AA minimum
Browser	Supported environments	Chrome 100+, Firefox 110+, Edge 100+, Safari 16+
Data	User deletion request processing	Within 30 days of request
PDF Output	Resume and cover letter rendering	Server-side; consistent across all supported browsers

11. Risks & Mitigations

ID	Risk	Impact	Likelihood	Mitigation
R-01	AI extracts inaccurate fields from complex job postings	High	Medium	Strict Gemini JSON schema enforced in prompt; Go validates all required fields before any write; user edits in side panel before commit
R-02	URL ingestion fails on login-walled or JavaScript-heavy pages	Low	High	Phase 1 treats URL ingestion as a best-effort convenience; Smart Paste is the primary path and always works; graceful fallback message on URL failure
R-03	GCP free tier quota exceeded under traffic	Low	Low	TanStack Query client caching reduces Firestore reads; Cloud Run scales to zero when idle; free tier limits are generous at Phase 1 volumes
R-04	Gemini API pricing or availability changes	Medium	Medium	AI call architecture is provider-agnostic at the service layer; switching models requires only a config change and prompt adjustment
R-05	Low Pro tier conversion reduces revenue viability	High	Medium	Free tier delivers standalone value; Pro features are meaningfully differentiated; match score and cover letter generation are demonstrably useful even in a trial
R-06	Resume PDF rendering inconsistency	Medium	Medium	Server-side PDF via chromedp headless Chrome; browser print-to-PDF not used; tested across all supported browsers

12. Phased Roadmap

Phase	Name	Timeline	Deliverables
Phase 1	MVP — Core Product	Weeks 1–12	Smart Paste + URL ingestion, side panel with confidence indicators, Kanban / Calendar / Table dashboard, status lifecycle with history, basic analytics (5 charts + 4 KPIs), user onboarding + profile with resume extraction, Pro page — resume canvas + AI suggestions + cover letter generator + PDF download, Free and Pro access gating
Phase 2	Automation Layer	Weeks 13–24	Browser extension (Chrome and Firefox) — one-click save from any job board, email webhook integration — auto-updates status from recruiter emails, AI match engine improvements — richer scoring model, enhanced ATS analysis
Phase 3	Copilot Layer	Weeks 25+	Interview preparation — AI-generated practice questions per role and company, predictive analytics — estimated response probability and time-to-decision, referral programme and community features

Sprint Plan — Phase 1

Sprint	Timeline	Deliverables
Sprint 1	Weeks 1–2	Go router, Firestore connection, Firebase Auth, React + Vite scaffold, authentication flows, route protection, CI pipeline
Sprint 2	Weeks 3–4	Smart Paste + URL ingestion endpoints, Gemini extraction integration, side panel with confidence indicators, field editing, commit flow
Sprint 3	Weeks 5–6	User onboarding 4-step wizard, resume upload and UserProfile extraction, profile page with editable fields
Sprint 4	Weeks 7–8	Kanban board with dnd-kit, Calendar view with FullCalendar, Table view with TanStack Table, status update pipeline and history log
Sprint 5	Weeks 9–10	Pro page — resume canvas, content editor with live sync, AI match score, missing skills, ATS risk, resume bullet suggestions, cover letter generator, PDF download
Sprint 6	Weeks 11–12	Analytics — 5 charts + 4 KPI cards with Recharts, date range filter, Free/Pro gating, empty states, loading skeletons, error boundaries, QA, accessibility audit, launch

Approval & Authorisation

Role	Name	Signature	Date
Product Owner			
Technical Lead — VybzTech			
Reviewed By			

TRACKD Product Requirements Document | VybzTech Inc. | Version 1.0 Draft | July 2026 | CONFIDENTIAL