TRACKD

Product Requirements Document — Two-Sided Edition

SCRAPER & INGESTION
CANDIDATE PIPELINE
AI INSIGHTS
RECRUITER HUB
SIMULATED INTEGRATION

PREPARED BY VybzTech Inc.

PRODUCT TRACKD — AI-Powered Two-Sided Job Platform

VERSION 2.0 | Draft for Stakeholder Review — supersedes v1.0

DATE July 2026

CLASSIFICATION Confidential

Table of Contents

1. Overview — What TRACKD Is, and What It Became

TRACKD started as a single-player tool: capture a job posting, structure it, track it through a pipeline, and use AI to make each application stronger. That product still exists, unchanged in ambition, inside this document. What changed is the shape of the company built around it.

TRACKD is now a two-sided platform. Candidates use it to eliminate the manual overhead of a job search — capturing opportunities with zero typing, tracking them in a real pipeline, and generating tailored application materials in minutes rather than hours. Recruiters use the same underlying data to triage applicants without drowning in generic, AI-inflated resumes. Both sides read and write the same structured record for a given opportunity — there is one data spine, not two disconnected products wearing the same logo.

This is a deliberate change from treating the recruiter side as a distant Phase 2. Building both sides together means the data model, the AI insight layer, and the mechanism that lets a recruiter act inside TRACKD without a real ATS integration (**Simulated Integration**, §9.4) are load-bearing from the first release — not features retrofitted onto a candidate-only schema later.

1.1 The Five Pillars

Five things make TRACKD what it is. The first four are inherited directly from the original candidate-only product; the fifth is new.

- **The Scraper** — gets job data in with zero manual typing, for the candidate.
- **The Dashboard** — gives the candidate a pipeline they can see and manage.
- **Analytics** — surfaces honest numbers on how a candidate’s search is actually going.
- **The AI Insight & Self-Improvement Layer** — tells a candidate where they are strong or exposed for a specific role, then gives them the tools to fix it.
- **The Recruiter Hub** — gives a recruiter a pre-scored, noise-free view of applicants, and a way to act on that view without needing their IT department’s approval first.

1.2 Document Scope

This is the single authoritative reference for what TRACKD must do in its v2.0 release — the first release where both sides ship together. It supersedes the v1.0 candidate-only PRD in full. Where a section is materially unchanged from v1.0, that is stated explicitly rather than left ambiguous.

2. The Problem We Are Solving

The v1.0 PRD made the candidate-side case in detail; it still holds. What follows restates it briefly and adds the recruiter-side half of the same broken exchange — because the two problems are, in fact, one problem viewed from opposite ends.

2.1 The Candidate Side

- **High-volume and manual.** A competitive job seeker applies to 30–150 roles per cycle, re-typing the same information and tailoring the same resume every time. Quality visibly drops after the fifteenth near-identical cover letter.
- **Data is nowhere and everywhere.** A spreadsheet abandoned three weeks in, a folder of screenshots, a wall of open tabs — no single place shows every application, its status, and what needs to happen next.
- **No performance data.** Candidates finish a search without knowing their response rate, their time-to-first-response, or which platforms actually convert. The only feedback loop left is intuition, and intuition says “try harder,” which is not useful advice.
- **Generic materials.** Writing a genuinely tailored letter for every role is time-prohibitive by hand, so the result reads like what it is — generic — and performs worse in both human and automated screening.

2.2 The Recruiter Side — new for v2.0

- **Signal-to-noise collapse.** Generative AI has made every inbound application look polished and keyword-optimized whether or not the candidate is actually qualified. Volume is no longer a proxy for interest, and keyword match is no longer a proxy for fit.
- **Manual triage cost.** A recruiter reading flat PDFs to find the handful worth a real look is spending hours per role on a task that structured data could compress to minutes.
- **New-tool fatigue.** Recruiters already live inside an ATS (Greenhouse, Lever, Workday). A standalone tool that adds a login without an immediate, obvious efficiency win is dead on arrival — this is precisely why real ATS integration is not a v1 requirement (§9.4).
- **No structured signal from candidates.** A flat resume tells a recruiter what a candidate claims. Nothing in a standard application structures that claim against the role’s actual requirements the way TRACKD’s match score and ATS-risk analysis already do for the candidate’s own use.

2.3 The Design Principle That Follows

Every TRACKD feature must reduce friction on the input side, improve visibility on the management side, or improve decision quality on the output side — for at least one of the two sides of the platform. Features that serve neither side, on either end, do not belong in this release.

3. Scope — What We Are Building, and What We Are Not

Both sides ship together in v2.0. The table below marks which side each capability belongs to and which phase it lands in — candidate-side rows are materially unchanged from v1.0; recruiter-side rows are new.

Feature	Side	Phase
Scraper — Smart Paste + URL ingestion	Candidate	Phase 1
Inbox + Explore (Sift scraper feed)	Candidate	Phase 1
Dashboard — Kanban, Calendar, Table	Candidate	Phase 1
Status tracking with full history log	Candidate	Phase 1
Analytics — 5 charts + 4 KPI cards	Candidate	Phase 1
AI match score, missing skills, ATS keyword risk	Candidate	Phase 1
Resume canvas + cover letter generator	Candidate	Phase 1
Companies — workspace + role posting	Recruiter	Phase 1
Applications — applicant tracking grid	Recruiter	Phase 1
Recruiter Hub — Candidate Storyline + fit scorecard	Recruiter	Phase 1
Simulated Integration action buttons	Recruiter	Phase 1
User profile + onboarding (both flows)	Both	Phase 1
Free / Pro / Recruiter tier access gates	Both	Phase 1
Browser extension — one-click save	Candidate	Phase 2
Real ATS webhook sync (Merge.to-class)	Recruiter	Phase 2
Email webhook — auto-status from recruiter emails	Candidate	Phase 2
Interview prep — AI practice questions	Candidate	Phase 3
Predictive analytics — response probability models	Both	Phase 3

3.1 Out of Scope for v2.0

- Native mobile app (iOS or Android)
- Real, OAuth-based ATS integration — covered by Simulated Integration until Phase 2
- Multi-seat recruiter team permissions / role hierarchies — v2.0 is single-admin per company
- Direct job application submission on the candidate’s behalf
- Third-party calendar sync (Google Calendar, Outlook)
- Automatic job-board scraping without a user- or recruiter-provided input

4. The Scraper — Two Ways In

Unchanged in behaviour from v1.0. Backend references updated: Claude (primary) with Gemini Flash (automatic fallback) replaces Gemini-only; Supabase replaces Firestore.

The single biggest source of drop-off in any tracking tool is the moment the user has to type. TRACKD removes that moment entirely, on the candidate side, through two ingestion methods that both end at the same place — a structured job record that populates the dashboard.

4.1 Method One — Smart Paste

The candidate copies a job description from wherever they found it and pastes it into the Smart Input Area. They click Parse. A Supabase Edge Function receives the raw text, strips residual formatting noise, and submits it to Claude with a strict JSON schema defining every field the system needs. If Claude errors, times out, or is rate-limited, the same request is automatically retried once against Gemini Flash with the identical schema — the candidate never sees the difference. The side panel opens once a valid, schema-conformant response is returned.

4.2 Method Two — URL Ingestion

The candidate pastes a URL from a public job listing. The edge function fetches the page, strips navigation and irrelevant elements, and extracts clean body text — from that point the flow is identical to Smart Paste. URL ingestion fails gracefully: a login wall, redirect, or 403 response returns a clear error prompting a switch to Smart Paste. The system never pretends a parse succeeded when it did not.

4.3 The Ingestion Side Panel

Every parse produces a side panel that slides in before anything is written to the database. Each field carries a confidence indicator — green, amber, or red — so the candidate knows exactly where to look before committing. The Commit button stays disabled until, at minimum, Company Name and Role Title are populated.

Note — shared schema with recruiter-side role posting

The same extraction schema now powers recruiter-side role posting (§9.1): a recruiter pasting a job description to post a role, and a candidate pasting the same listing to track it, both produce data in the identical shape. This is what makes the two sides of the platform one data spine rather than a coincidence of naming.

5. Job Tracking — The Dashboard

Unchanged from v1.0.

Once a job is committed, it lands in the dashboard — the candidate's central workspace for managing their pipeline day to day. Three views of the same underlying data; the candidate switches freely and the data stays in sync.

5.1 Kanban View

The default view. Five columns represent the five possible states of any application. Cards move between columns by drag-and-drop; every move is a status update, logged to the job's history. Each card shows company name and logo, job title, date logged, source, and — for Pro users — the AI match score as a small circular badge.

Column	Status Value	Meaning
Saved	Bookmarked	Captured but not yet applied to
Applied	Applied	Application has been submitted
Interviewing	Interviewing	At least one interview confirmed or completed
Offer Received	Offer	A formal offer has been made
Rejected / Closed	Rejected	Declined, or the role has been filled

5.2 Calendar View

Jobs appear on their applied date; a second marker appears on any captured application deadline. Month, week, and day views are supported. Solves the deadline blind-spot — the candidate sees at a glance what is coming up.

5.3 Table View

A dense, sortable, filterable grid. Status and source are multi-select filterable; status is editable inline. Row checkboxes enable bulk archive or delete.

5.4 Status Lifecycle

Status transitions follow a defined matrix — a candidate cannot move directly from Saved to Offer, for example. Every transition is recorded with a timestamp, and this history powers the time-in-stage analytics. The same matrix, encoded once in **statusMachine.ts**, also governs the recruiter-side Simulated Integration actions (§9.4) — a recruiter's "Move to Interview" click is the same state transition as the candidate dragging their own card.

From Status	Allowed Next Statuses
Saved (bookmarked)	Applied, Rejected
Applied	Interviewing, Rejected
Interviewing	Offer Received, Rejected
Offer Received	Terminal — no further transition

From Status	Allowed Next Statuses
Rejected	Terminal — no further transition

6. Analytics

Unchanged from v1.0.

The analytics page exists for one reason: honest data about the candidate's own search. Not encouragement, not gamification — data.

6.1 KPI Cards

Metric	Calculation	Why It Matters
Total Applications	Count of all non-archived tracked jobs	Sense of volume and effort invested
Response Rate	(Interviewing + Offer) / Applied × 100%	Honest measure of pipeline conversion
Offer Rate	Offer / Applied × 100%	The number that matters most at the end
Avg. Time to Response	Mean days, Applied → first Interviewing entry	Identifies where the delay actually is

6.2 Chart Suite

Five charts, with a date-range filter (Last 30 Days / Last 90 Days / All Time) applying to all of them simultaneously: Applications Over Time (area line), Pipeline Breakdown (donut), Applications by Source (horizontal bar), Top Roles Applied To (ranked list), and Average Time in Stage (stat cards per status).

Note on data quality

Analytics are only as good as the candidate's own status-update discipline. If updates stop, analytics go stale — the product shows an honest warning rather than a false picture.

7. AI Insights

Behaviour unchanged from v1.0. Vendor updated: Claude (primary) with Gemini Flash (automatic fallback) replaces Gemini-only.

The AI insight layer is what separates TRACKD from a well-designed spreadsheet. It takes the candidate’s profile and cross-references it against each specific job they have tracked. The output is a per-job diagnosis, not a generic score.

7.1 Match Score

A 0–100 percentage shown as a circular gauge — green at 70+, amber 40–69, red below 40. Calculated by comparing the candidate’s extracted skills, experience, and seniority against the job’s stated requirements. TRACKD is explicit with the candidate that this is a directional signal, not a guarantee — **a tool, not a verdict**.

7.2 ATS Keyword Analysis

A risk badge (Low / Medium / High) based on how many required keywords are absent from the candidate’s resume, an expandable list of the specific flagged keywords, and where each one appeared in the job description. The loop between insight and action closes on the same page — accepting a suggestion in the resume editor updates the risk level live.

7.3 Resume Bullet Suggestions

The most opinionated part of the AI layer: rewritten action verbs, quantified impact, and job-matched terminology — without misrepresenting the candidate’s actual experience. Nothing is automatic; the candidate accepts or dismisses every suggestion individually, and the canvas updates immediately on acceptance.

8. Self-Improvement — The Empowerment Layer

Unchanged from v1.0. PDF generation updated: see §11.

The AI Insights panel tells the candidate what is wrong. This layer gives them the tools to fix it without leaving TRACKD.

8.1 The Resume Canvas

A split-pane layout: left pane renders a print-ready, live resume; right pane is the tabbed content editor (Summary, Experience, Skills, Education). Every edit updates the left pane in under 200 milliseconds. Download produces a server-generated PDF, named *FirstName_LastName_CompanyName_Resume.pdf*.

8.2 Cover Letter Generator

The candidate selects a tone — Professional, Confident, or Conversational — and Claude produces a full letter using the job description and profile as context. Fully editable before download or copy; Regenerate is available with a confirmation prompt.

8.3 The Growth Loop

Analytics shows a low response rate. AI Insights shows high ATS-keyword risk across recent applications. The Resume Canvas lets the candidate fix it. They apply again, and wait for the next data point. Over multiple iterations, match scores and response rates genuinely improve — the product has an active interest in the candidate's success, not just their continued engagement.

9. The Recruiter Side — New for v2.0

This section did not exist in v1.0. It is the product-level record of the two-sided pivot.

Where Sections 4–8 describe the candidate's half of TRACKD, this section describes the recruiter's: a workspace for posting roles, triaging applicants, and acting on that triage — without needing a real ATS integration to get value on day one.

9.1 Companies

A Company is a recruiter's workspace — the container for their posted roles and, later, their team. Creating one takes a name, industry, and size; nothing more is required to start posting. Posting a role reuses the exact extraction schema from candidate-side ingestion (§4) — a recruiter pastes or links a job description, Claude structures it, and the recruiter reviews the same kind of confidence-scored side panel a candidate sees. This is a deliberate reuse, not a coincidence: it means the ingestion UI and its edge function are built once and used by both sides.

Component	Behaviour
Role posting	Title, description, department, seniority, compensation band — extracted via the shared schema
Role list	Applicant count, days open, status (Open / Paused / Filled) per role
Team (v2.0)	Single admin per company — structured now so multi-seat is a UI change in Phase 2, not a schema change

9.2 Applications — Applicant Tracking Grid

For a given role, the recruiter sees every candidate who has applied through TRACKD, in an organized, searchable grid — not a stack of individually-opened flat resumes. Each row shows the AI match score, current status, source, and applied date. The grid is filterable by score range, status, source, and seniority, and searchable by name, skill, or resume keyword.

9.3 Recruiter Hub — The Candidate Storyline

Opening a candidate from the Applications grid opens the Recruiter Hub — an enterprise hiring-desk framework built around the **Candidate Storyline**: a visual summary of skills, experience arc, and stated fit for the specific role, generated from the same AI layer that powers the candidate's own Pro page. This is the direct answer to the recruiter's actual question — not "does this resume contain the right words," but "does this candidate genuinely fit, and where would they need to grow into the role." A flat PDF cannot answer that; a structured, AI-cross-referenced record can.

9.4 Simulated Integration — The Adoption Mechanic

This is the single most important new idea in v2.0, and it deserves to be stated plainly rather than buried in a feature table.

Real ATS integration — OAuth into Greenhouse, Lever, Workday — requires enterprise IT approval and a security review. That is a multi-month sales cycle standing between a recruiter liking TRACKD and actually being allowed to use it. **Simulated Integration** removes that barrier for v2.0: a recruiter clicks an action button —

Shortlist, Move to Interview, Reject — inside the TRACKD Recruiter Hub, and the platform instantly updates the candidate's own Kanban board and sends them a transactional status email. No real webhook to a third-party ATS exists underneath this action in v2.0.

To the candidate, this feels like the employer has a sophisticated, integrated tracking system — because functionally, from their side, it behaves like one. To the recruiter, it is a one-click, zero-IT-ticket workflow that produces the same dopamine hit as a real integration, without the months of approval it would otherwise require. TRACKD is honest about this mechanism in its own product copy: it is never marketed as a real ATS connection, because it is not one — the value (speed, no approval process) is real even though the underlying integration is not, and that distinction has to survive contact with legal review (see the companion Legal Framework document) as much as it survives contact with the UI.

Once a recruiter is genuinely hooked on the workflow, Phase 2 introduces the real thing — Unified API-class integrations (Merge.to / Unified.to) so a status change in Greenhouse pushes a webhook into TRACKD automatically, and vice versa. Simulated Integration is the bridge to that point, not a permanent substitute for it.

10. User Stories & Acceptance Criteria

10.1 Candidate

ID	Story	Acceptance Criteria
US-C01	Paste a job description and have all fields auto-populated	Side panel opens within 8s of clicking Parse, all extractable fields populated
US-C02	Correct AI-extracted fields before committing	Every field is editable; changes persist to Supabase on commit
US-C03	Drag a job card to update its status	Status updates and history logs within 2s; all views reflect it
US-C04	See a match score for each job (Pro)	Match score, missing skills, ATS risk shown within 8s of Pro page load
US-C05	Accept an AI resume suggestion	Editor and canvas update within 200ms
US-C06	Download a tailored resume PDF	Correctly formatted PDF downloads within 5s
US-C07	Generate a cover letter with a chosen tone	Relevant letter appears in the output area within 8s
US-C08	See charts of my application activity	Given 5+ applications, all 5 charts and 4 KPI cards show correct values

10.2 Recruiter — new for v2.0

ID	Story	Acceptance Criteria
US-R01	Post a role and see it live in my Company workspace	Role appears in the role list immediately after extraction confirms
US-R02	View all applicants to a role in one grid	Grid loads within 3s for up to 200 applicants; sortable and filterable
US-R03	See why a candidate fits without reading a full resume	Storyline and fit scorecard render within 5s of opening a candidate
US-R04	Move a candidate to Interview with one click	Candidate’s status updates within 2s; email sent within 1 minute
US-R05	Search candidates by skill across all applicants	Results filter live while typing, no full page reload

11. Non-Functional Requirements

Category	Requirement	Target
Performance	Ingestion response (excl. AI latency)	< 3 seconds
Performance	Claude/Gemini extraction response	< 8 seconds (P95)
Performance	Dashboard load (50 jobs)	< 3 seconds (P95)
Performance	Applications grid load (200 applicants)	< 3 seconds (P95)
Performance	Resume canvas edit-to-render	< 200ms
Availability	Platform uptime	99.5% monthly SLA
Scalability	Concurrent users without degradation	1,000+ via Supabase + edge function auto-scaling
Security	Authentication	Supabase Auth (GoTrue), JWT + Row Level Security on every table
Security	Data encryption	Postgres at-rest encryption; HTTPS enforced
Accessibility	Standard	WCAG 2.1 AA minimum, both candidate and recruiter UI
Browser	Supported environments	Chrome 100+, Firefox 110+, Edge 100+, Safari 16+
Data	User deletion request processing	Within 30 days of request, both account types
PDF Output	Resume and cover letter rendering	@react-pdf/renderer in a Supabase Edge Function; consistent across all supported browsers — see the companion Backend Doc §7 for the full resolution and fallback plan

12. Risks & Mitigations

ID	Risk	Impact / Likelihood	Mitigation
R01	AI extracts inaccurate fields from complex postings	High / Medium	Strict Claude/Gemini JSON schema; edge function validates before any write; candidate/recruiter edits before commit
R02	Building both sides at once slips the timeline	High / High	Shared data model and shared ingestion schema mean recruiter-side work reuses candidate-side code rather than duplicating it
R03	Recruiters see zero candidates on day one and bounce	High / High	Seed the Recruiter Hub with realistic sample candidates so day-one value is visible before organic volume exists
R04	Simulated Integration perceived as “fake” once understood	Medium / Medium	Never claim real ATS integration in product copy or marketing; the speed and no-IT-ticket value is real regardless
R05	Claude/Gemini pricing or availability shifts	Medium / Medium	Provider-agnostic internal call interface; automatic fallback already built, not a future change
R06	Supabase RLS misconfiguration exposes cross-tenant data	High / Low	Per-table RLS policy review before launch; recruiter access scoped strictly to their own company’s applicants
R07	Low Pro / Recruiter tier conversion reduces revenue viability	High / Medium	Free tier delivers standalone value on both sides; Recruiter tier pricing deferred to a dedicated workshop rather than guessed

13. Phased Roadmap

Phase	Name	Deliverables
Phase 1	Two-Sided MVP	Candidate: Smart Paste + URL ingestion, Kanban/Calendar/Table, analytics, Pro page (resume canvas, match score, cover letters). Recruiter: Companies, Applications grid, Recruiter Hub, Simulated Integration. Shared: onboarding, Free/Pro/Recruiter gating.
Phase 2	Real Integration Layer	Browser extension (candidate), real ATS webhook sync via Merge.to-class integration (recruiter), email-triggered status updates, richer AI match model, multi-seat recruiter teams.
Phase 3	Copilot Layer	Interview preparation (candidate), predictive analytics (both sides), referral programme and community features.

Within Phase 1 itself, candidate-side and recruiter-side workstreams run in parallel rather than sequentially — see the companion Architecture & Feature List and Build UI Prompt documents for the module-by-module build order that makes this practical without one side blocking the other.

14. Approval & Authorisation

Role	Name	Signature	Date
Product Owner			
Technical Lead — VybzTech			
Reviewed By			

TRACKD Product Requirements Document | VybzTech Inc. | Version 2.0 — Two-Sided Edition | July 2026 | CONFIDENTIAL | Supersedes v1.0 in full